

# A brief introduction to Python & Numpy

*Deep Learning for Computer Vision*

TA: Junming CHEN

COMP 4471 & ELEC4240 Lab 1  
September 14



# About Myself

---

- Name: Junming CHEN (Jeremy)
- Third year AI PhD student
- Research interests:
  - Video understanding & synthesis
  - Multi-modal learning & 3D vision
  - Domain adaptation & generalization
- Email
  - [jeremy.junming.chen@connect.ust.hk](mailto:jeremy.junming.chen@connect.ust.hk)



# Today's Tutorial

---

- Introduction to Python
  - Containers
  - Functions
  - Classes
- Introduction to Numpy
  - Basics
  - Useful functions
  - Common pitfalls

# An Introduction to Python

---

- What is Python?
  - Python is a high-level, dynamically typed multiparadigm programming language.
- Why do we learn Python?
  - Simplicity: Natural language style
  - Huge Community: Libraries and Frameworks
  - Most popular language used by AI researchers
- Now start your favorite python terminal
  - Recommend: Vscode, Jupyter Notebook
  - Environment management: [Anaconda is all you need!](#)

# Python Basics

---

- Basic data types
  - Numbers: int, float, etc.
  - Booleans: True, False
    - and, or, not, !=, ==
  - Strings

```
In []: hw = 'hello' + ' world!'
```

- `replace()`, `split()`, etc.

# Python Containers

- Lists: ordered element sequence

```
In []: x = [1,2,3]
```

- List could contain different types of elements

```
In []: x[0] = 'hello'
```

- Basic functions

```
In []: x.append(4.0)
```

- Loops

```
In []: for element in x:  
        print(element)
```

# Python Containers

- Dictionaries: (key, value) pairs

```
In []: d = {'cat': 'cute', 'dog': 'furry'}
```

- Entry to value

```
In []: print(d['cat'])
```

- Set an entry

```
In []: d['fish'] = 'wet'
```

- Loops

```
In []: for animal, adj in d.items():  
        print('A %s is %d' % (animal, adj))
```

# Python Containers

- **Sets:** An unordered collection of distinct elements

```
In []: animals = {'cat', 'dog'}
```

- Add / remove elements to set

```
In []: animals.add('fish')  
In []: animals.remove('cat')
```

- **Tuple:** An (immutable) ordered list of values

```
In []: d = {(x, x + 1): x for x in range(10)}
```



# Python functions

- Define and call

```
def sign(x):  
    if x > 0:  
        return 'positive'  
    elif x < 0:  
        return 'negative'  
    else:  
        return 'zero'  
  
for x in [-1, 0, 1]:  
    print(sign(x))
```

- Default argument

```
def hello(name, Loud=False):  
    if loud:  
        print('HELLO, %s!' % name.upper())  
    else:  
        print('Hello, %s' % name)  
  
hello('Bob') # Prints "Hello, Bob"  
hello('Fred', Loud=True) # Prints "HELLO, FRED!"
```

# Python classes

- Define

```
class Greeter(object):  
  
    # Constructor  
    def __init__(self, name):  
        self.name = name # Create an instance variable  
  
    # Instance method  
    def greet(self, loud=False):  
        if loud:  
            print('HELLO, %s!' % self.name.upper())  
        else:  
            print('Hello, %s' % self.name)
```

- Construct an instance

```
g = Greeter('Fred') # Construct an instance of the Greeter class  
g.greet()           # Call an instance method; prints "Hello, Fred"  
g.greet(loud=True)  # Call an instance method; prints "HELLO, FRED!"
```

# An Introduction to Numpy

---

- What is Numpy?
  - A scientific computing library for Python
  - All you need for **Linear Algebra on CPU**
  - Want linear algebra on GPU? PyTorch / TensorFlow / JAX
- Why do we learn Numpy?
  - N-dimensional array manipulation
    - E.g., your dataset (images), weights, and feature maps...
  - Deep learning is based on linear algebra and optimization

# Numpy Basics

---

- Importing

```
In []: import numpy as np
```

- Array creation

```
In []: a = np.array([1, 2, 3])
```

- Shape: size of the array

```
In []: print(a.shape)
```

# Array Creation

- Array creation (a 2 by 3 matrix)

```
In []: b = np.array([[1,2,3],[4,5,6]])
```

- Shape: size of the array

```
In []: print(b.shape)
```

- Specify data type:

```
In []: b = np.array([[1,2,3],[4,5,6]],  
dtype=np.int64)  
In []: print(b.dtype)
```

# Array creation: useful functions

---

- `np.zeros()`
- `np.ones()`
- `np.full()`
- `np.eyes()`
- `np.random.random()`
- Specify the shape with Python tuple

# Array Indexing

- Fetching one number in the n-D array

```
In []: a = np.array([1, 2, 3])
```

- Use square brackets

```
In []: print(a[0], a[1], a[2])
```

# Array Indexing

- Fetching one number in the n-D array

```
In []: b = np.array([[1,2,3],[4,5,6]])
```

- Use square brackets

```
In []: print(b[0, 0], b[0, 1],  
             b[1, 0])
```



# Array Indexing

- Fetching a sub-array of the n-D array

```
In []: a = np.array([[1,2,3,4],  
                    [5,6,7,8], [9,10,11,12]])
```

- Just like that for Python lists

```
In []: b = a[:2, 1:3]
```

- **Warning:**
  - Index starts at 0
  - Last index is not included
  - 1-2 rows, 2-3 columns

# Array Indexing

- Fetching a sub-array of the n-D array

```
In []: a = np.array([[1,2,3,4],  
                    [5,6,7,8], [9,10,11,12]])
```

- Just like that for Python lists

```
In []: b = a[:2, 1:3]
```

- **Warning:**
  - Slicing does not copy arrays
  - `b[0, 0] = 77` change `a[0, 1]`

# Integer Indexing

- Want to index many elements

```
In []: print(np.array(  
    [a[0, 0], a[1, 1],  
    a[2, 0], a[2, 1]]))
```

- Concise way: use less brackets

```
In []: print(a[[0, 1, 2, 2],  
               [0, 1, 0, 1]])
```

# Boolean Indexing

---

- Select elements that satisfy certain conditions

```
In []: a = np.array([[1,2], [3, 4],  
[5, 6]])
```

```
In []: bool_idx = (a > 2)
```

```
In []: print(a[bool_idx])
```

# Boolean Indexing

---

- Select elements that satisfy certain conditions

```
In []: a = np.array([[1,2], [3, 4],  
[5, 6]])
```

```
In []: bool_idx = (np.sqrt(a) > 2)
```

```
In []: print(a[bool_idx])
```

# Array arithmetic

---

- Assume  $x$  and  $y$  have the same shape

```
In []: print(x + y)
```

```
In []: print(x - y)
```

```
In []: print(x*y)
```

```
In []: print(x/y)
```

```
In []: print(np.sqrt(x))
```

# Array arithmetic

- How to do matrix multiplication?

```
In []: print(np.dot(x, y))
```

```
In []: print(x.dot(y)) # Simpler: x@y
```

- Transpose

```
In []: x = np.array([[1,2], [3,4]]))
```

```
In []: print(x.T)
```

# Broadcasting

- Assume we want to deduct a scalar *mean* value for each row element

```
In []: a = np.array([[1, 2, 3],  
                    [4, 5, 6]])
```

```
In []: for i in range(a.shape[0]):  
        a[i, :] = a[i, :] - mean
```

- The concise (and faster) way

```
In []: a = a - mean
```



# Broadcasting

- Assume we want to deduct a three-dimension *mean* vector for each row

```
In []: a = np.array([[1, 2, 3],  
                    [4, 5, 6]])
```

```
In []: mean = np.sum(a, axis = 0)/2  
       print(mean)
```

```
Out []: [2, 3.5, 4.5]
```

# Broadcasting

- Assume we want to deduct a three-dimension *mean* vector for each row

```
In []: a = np.array([[1, 2, 3],  
                    [4, 5, 6]])
```

```
In []: for i in range(a.shape[0]):  
        a[i, :] = a[i, :] - mean
```

- The concise (and faster) way

```
In []: a = a - mean
```

# Broadcasting: how it works?

- Basically we are dealing with two arrays with different shapes

```
In []: a = np.array([[1, 2, 3],  
                    [4, 5, 6]])
```

```
In []: a = a - mean
```

↓       ↓  
(2, 3)   (3,)

# Broadcasting: how it works?

- Behind Idea: Broadcast one row and do element-wise subtraction

In `[]`: `a = a - mean`

↓      ↙  
(2, 3) (3,)

1	2	3
4	5	6

-

2	3.5	4.5
---	-----	-----

# Broadcasting: how it works?

- Broadcast one row and do element-wise subtraction

In `[]`: `a = a - mean`

↓      ↙  
(2, 3) (3,)

1	2	3
4	5	6

-

2	3.5	4.5
2	3.5	4.5




Broadcast

# Broadcasting: general rules

- When the two arrays have different shape, append 1 to the smaller array

In []:  $a = a - \text{mean}$

$\downarrow$        $\swarrow$   
(2, 3)    (3,)



(1, 3)

# Broadcasting: general rules

- Only two cases are compatible in broadcasting
  - They have the same size
  - One of the size is 1

In `[]`:  $a = a - \text{mean}$

$\downarrow$        $\downarrow$   
(2, 3)   (1, 3)

# Broadcasting: general rules

- Only two cases are allowed for each dimension:
  - They have the same size
  - One of the size is 1
- Broadcasting brings each dimension to be the maximum possible size

```
In []: a = a - mean
```

↓            ↓  
(2, 3)    (2, 3)




# Broadcasting: general rules

- When operating on two arrays, NumPy compares their shapes element-wise. Two dimensions are compatible when:
  - They are equal, or
  - One of them is 1
- Broadcasting brings each dimension to be the maximum possible size
- Exercise: add arrays of shapes (2, 3, 4) and (3, 4)
  - > (1, 3, 4) -> (2, 3, 4)

# Broadcasting: general rules

- More exercises: whether those array shapes are compatible for broadcasting
  - (3,6,8)  
(6,8)
  - (3,6,8)  
(3,6)
  - (3,6,8)  
(3,6,1)
  - (3,1,8)  
(1,6,1)

# Use broadcast



$$\text{dist}(X, Y) = \sqrt{(X - Y)^2} = \sqrt{X^2 + Y^2 - 2XY}$$

# Other packages

---

- Scipy
  - Optimization, eigenvalue problems, algebraic equations, differential equations
- PIL / OpenCV
  - Image manipulation
- Matplotlib
  - For data visualization
- **Documentation** and **Google** are your friends



# Q&A