

Intro to Tensorflow (TF)

Deep Learning for Computer Vision
ELEC4240/COMP4471

About me

- Name: Chenyang QI
- Email: cqiaa@connect.ust.hk
- 3rd year PhD student in CSE,
supervised by Prof. Qifeng Chen
- Mainly doing research in computational photography

Today's Tutorial

1. Quick overview of Tensorflow
2. Installation
3. Example {from loading data to evaluating model}
4. Save & Load model
5. Try it yourself

Source of tutorial : [tf_basic](#)

Read the documentation to understand more. This tutorial only provide **the basic**

What is tensorflow?

Open Source software library (by google)

Well known usage : machine learning; but has [other usage](#) as well



Why tensorflow?

1. You're already familiar with tensorflow
2. You work with google
3. You want faster deployment and production

In general :

More flexible than pytorch → more complicated

[MXNET](#) provides alternative but community is not as big

Installation

Latest version : Tensorflow 2.0

Command :


Pip install tensorflow

Conda install tensorflow

General guideline

1. Design your application {supervised/unsupervised}
2. Make sure you have necessary components (i.e. opensource, data&label)
3. Decide on the input & output → training & deployment have to be **consistent**
4. Test, fine tune, & debug (complexity, overfitting, etc)


Example - MNIST classification



Task :

Build a **model** to classify input image as digits

Example - MNIST classification



1. Get the data & label for training → preprocessing helps {observe}
2. Set the input & output for your model → based on how you want to deploy
3. Debug → until satisfactory

Quick overview

Implementation using Tensorflow - just **few lines**

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Deeper dive

Step1

Call the necessary package (**import**)

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step1

Call the necessary package (**import**)

Step2

Load the dataset (MNIST is **common** hence simple)

Preprocess data {normalize} → why? Does it matter?

Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step 1

Call the necessary package (**import**)

Step 2

Load the dataset (MNIST is **common** hence simple)

Preprocess data {normalize} → why? Does it matter?

Step 3

Design your model.

Is this model effective? Why this model?

What is the **input** and **output**?

Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step 1

Call the necessary package (**import**)

Step 2

Load the dataset (MNIST is **common** hence simple)

Preprocess data {normalize} → why? Does it matter?

Step 3

Design your model.

Is this model effective? Why this model?

What is the **input** and **output**?

Given images, what will be the output of model?

Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step 1

Call the necessary package (**import**)

Step 2

Load the dataset (MNIST is **common** hence simple)


Preprocess data {normalize} → why? Does it matter?

Step 3

Design your model.

Is this model effective? Why this model?

What is the **input** and **output**?



Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step1

Call the necessary package (**import**)

Step2

Load the dataset (MNIST is **common** hence simple)

Preprocess data {normalize} → why? Does it matter?

Step3

Design your model.

Is this model effective? Why this model?

What is the **input** and **output**?

```
predictions[0]
```

```
array([1.6707758e-06, 8.3274145e-08, 9.8423456e-08, 1.9251273e-07,
       1.4543222e-06, 2.4620399e-02, 8.9157339e-07, 4.9053874e-02,
       6.1236402e-05, 9.2625999e-01], dtype=float32)
```

```
np.argmax(predictions[0])
```

9

Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step 1

Call the necessary package (**import**)

Step 2

Load the dataset (MNIST is **common** hence simple)

Preprocess data {normalize} → why? Does it matter?

Step 3

Design your model.

Is this model effective? Why this model?

What is the **input** and **output**?

Step 4

Decide on how you want to train the model

Sparse_categorical vs categorical?

Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step 1

Call the necessary package (**import**)

Step 2

Load the dataset (MNIST is **common** hence simple)

Preprocess data {normalize} → why? Does it matter?

Step 3

Design your model.

Is this model effective? Why this model?

What is the **input** and **output**?

Step 4

Decide on how you want to train the model

Sparse_categorical vs categorical?

Sparse_categorical, label = [0, 2,]

categorical, label = [[1 0 0 0], [0 0 1 0],]

Deeper dive

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

Step1

Call the necessary package (**import**)

Step2

Load the dataset (MNIST is **common** hence simple)

Preprocess data {normalize} → why? Does it matter?

Step3

Design your model.

Is this model effective? Why this model?

What is the **input** and **output**?

Step4

Decide on how you want to train the model

Step5

Train & Evaluate on non-training set. (**why?**)

Model

Training is often expensive. After training, what to do?

- A. Close program & do nothing
- B. Save the training weight & model architecture

Model : Check

Debug & Check if model is correctly created:

```
# Display the model's architecture  
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130

Total params: 407,050

Trainable params: 407,050

Non-trainable params: 0

Model : Save

What is the difference?

```
# Save the weights  
model.save_weights('./checkpoints/my_checkpoint')
```

```
model.save('my_model.h5')
```

Model : Save & Load

What is the difference?

```
# Save the weights  
model.save_weights('./checkpoints/my_checkpoint')
```

```
model.save('my_model.h5')
```

```
# Restore the weights  
model.load_weights('./checkpoints/my_checkpoint')
```

Model : Save & Load

What is the output?

```
# Save the weights
model.save_weights('./checkpoints/my_checkpoint')
```

```
# Restore the weights
model.load_weights('./checkpoints/my_checkpoint')
```

```
def create_model():
    model = tf.keras.models.Sequential([
        keras.layers.Dense(512, activation='relu', input_shape=(784,)),
        keras.layers.Dropout(0.2),
        keras.layers.Dense(10)
    ])

def create_model_2():
    model = tf.keras.models.Sequential([
        keras.layers.Dense(512, activation='relu', input_shape=(784,)),
        keras.layers.Dropout(0.2),
        keras.layers.Dense(10)
    ])

model = create_model()
model_2 = create_model_2()

model1.save_weights(path)
model_2.load_weight(path)
```


Model : Save & Load

What is the output?

```
# Save the weights
model.save_weights('./checkpoints/my_checkpoint')
```

```
# Restore the weights
model.load_weights('./checkpoints/my_checkpoint')
```

```
def create_model():
    model = tf.keras.models.Sequential([
        keras.layers.Dense(512, activation='relu', input_shape=(784,)),
        keras.layers.Dropout(0.2),
        keras.layers.Dense(10)
    ])

def create_model_2():
    model = tf.keras.models.Sequential([
        keras.layers.Dense(600, activation='relu', input_shape=(784,)),
        keras.layers.Dropout(0.2),
        keras.layers.Dense(10)
    ])

model = create_model()
model_2 = create_model_2()

model1.save_weights(path)
model_2.load_weights(path)
```

Model : Save & Load

How about this one?

```
def create_model():  
    model = tf.keras.models.Sequential([  
        keras.layers.Dense(512, activation='relu', input_shape=(784,)),  
        keras.layers.Dropout(0.2),  
        keras.layers.Dense(10)  
    ])  
  
def create_model_2():  
    model = tf.keras.models.Sequential([  
        keras.layers.Dense(512, activation='relu', input_shape=(784,)),  
        keras.layers.Dropout(0.2),  
        keras.layers.Dense(10)  
    ])  
  
model = create_model()  
model_2 = create_model_2()  
  
model.save(path)  
model_2.load(path)
```

Model : Save & Load

Last one

```
def create_model():  
    model = tf.keras.models.Sequential([  
        keras.layers.Dense(512, activation='relu', input_shape=(784,)),  
        keras.layers.Dropout(0.2),  
        keras.layers.Dense(10)  
    ])  
  
def create_model_2():  
    model = tf.keras.models.Sequential([  
        keras.layers.Dense(512, activation='relu', input_shape=(784,)),  
        keras.layers.Dropout(0.2),  
        keras.layers.Dense(10)  
    ])  
  
model = create_model()  
model_2 = create_model_2()  
  
model.save(path)  
new_model = tf.keras.models.load_model(path)
```

Recap

Utilize model [saving & loading](#) → checkpoint at every few iters

Try it yourself : Fashion MNIST

Check the jupyter notebook file and try to build your own model → self explanatory

Do not trust the code as is. I modified the code [hint: related to tutorial]

Take knowledge from today's tutorial and try to build a classifier


But before that,

p.s. → might **be easier** to write your own code from scratch

→ [solution here](#)

Another simple dataset : Fashion MNIST

It's just like MNIST but with clothing




Label	Class
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

Might be helpful for debugging

1. Check input → is it correct?


```
plt.figure()  
plt.imshow(train_images[0])  
plt.colorbar()  
plt.grid(False)  
plt.show()
```




Might be helpful for debugging

1. Check input → is it correct?
2. Verify predictions

```
i = 0
plt.figure(figsize=(6,3))
plt.subplot(1,2,1)
plot_image(i, predictions[i], test_labels, test_images)
plt.subplot(1,2,2)
plot_value_array(i, predictions[i], test_labels)
plt.show()
```



Ankle boot 100% (Ankle boot)



Have fun!